WHERE WE ARE / WHAT IS LEFT / WHAT TO DO NEXT

PRAMAAN hackathon war room

Offline-first forensic evidence analysis platform — audited against the working tree on 04 Sep 2026 by running the suites, not by reading the code.

> **BOTTOM LINE** The backend is done and green. But **9,725 lines of the frontend redesign are dark code**: they score nothing, and they break the build. Close the build first — everything else queues behind it.

VERIFIED TODAY — WHAT IS DONE

Backend	14,169 LOC application + 8,738 LOC tests. Full pytest suite green (exit 0).
Forensic pipeline	Ingest, SHA-256, pHash / dHash / aHash, EXIF and ISO-BMFF metadata, exact perceptual index, near-duplicate retrieval, propagation and earliest known instance, C2PA provenance, compression forensics, transparent weighted fusion, hash-chained audit log, PDF report. ~50 endpoints.
Detector engine	2,264 LOC, **57 / 57 tests**. Digest-verified real weights: image Swin-B 347 MB, audio Wav2Vec2-large 1.26 GB.
API contract	Recorded from the live app and replayed by the frontend: **72 / 72 checks pass**.
Deployed	Render backend + Vercel frontend, CORS and the `/api/backend/*` proxy wired and working.
Live console	10 screens, 12,733 LOC committed — this is what a judge sees today.

THE RUBRIC WE ARE SCORED ON

40%	**Functionality Completion**	Objectives met, functional requirements fulfilled, effectiveness against problem complexity
40%	**Quality of Deliverable**	Code quality, documentation, scalability, reusability, maintainability
10%	**User Experience and Design**	Intuitive, aesthetic, navigable, inclusive, accessible
10%	**Presentation and Communication**	Problem, solution, impact, storytelling, value proposition

Plus a complexity multiplier: problem statements are tiered High / Medium / Low, and depth on a harder statement is weighted up. **Read: 80 percent of the score is code and completeness.** UI polish is capped at 10 percent — and a half-migrated tree with two parallel CSS systems actively costs us on the 40 percent that matters most.

THE REAL GAPS

What is left, ranked by what it costs us

BLOCKER	**The redesign is dark code.** 9,725 LOC — 9 components, 4 libs, 6 CSS files — import only each other. No screen, no `App.tsx`, no `main.tsx` references them.	10% UX, drags 40% quality
BLOCKER	**The working tree does not build.** 16 TypeScript errors: `IconName` is missing 14 glyphs the new components use, and `Media.tsx` imports a non-existent `mediaIcon`. `tsc -b` fails, so `npm run build` fails, so nothing can deploy.	40% functionality
BLOCKER	**Two CSS systems, neither complete.** `main.tsx` loads `global.css` (old, 3,520 lines) and never `styles/index.css` (new, 5,748 lines). Of 532 classes used in TSX: 291 exist only in the new CSS, 159 only in the old, and **40 are defined nowhere** — the print/report view and the propagation graph are broken right now.	40% quality, 10% UX
HIGH	**The production AI detector is switched off.** `PRAMAAN_ENABLE_AI_DETECTOR=false` in `render.yaml` (free-tier RAM). A judge who opens the live URL watches the flagship deepfake signal abstain.	40% functionality
MEDIUM	Four orphan modules nothing imports: `Artefact.tsx`, `Async.tsx`, `Overlays.tsx`, `lib/caseflow.ts`.	40% quality
MEDIUM	No `video_detector.pt`, so video is honestly UNAVAILABLE — defensible, but a visible hole in a multi-modal claim.	40% functionality
MEDIUM	**No measured accuracy published anywhere.** `scripts/test_phash.py` exists but its report reaches neither the README nor the deck. Thresholds and weights are uncalibrated prototype defaults.	40% functionality
MEDIUM	Fusion weights in the workflow spec disagree with the implemented weights. Unresolved product decision sitting in the middle of the scoring engine.	40% quality
LOW	Docker image never actually built or run; README test count is stale; the deck is from 25 Aug and predates everything since.	40% quality, 10% presentation

DO NEXT, IN THIS ORDER

1. **Close the icon set** ~1h
   Add the 14 missing names and SVG paths to `Icon.tsx`, export `mediaIcon`. Done when `npm run typecheck` is clean.
2. **Switch the stylesheet entry** ~30m
   `main.tsx` to `./styles/index.css`. Keep `global.css` alongside for one step and confirm no screen regresses.
3. **Finish the migration** ~3-4h
   Move the screens off the 159 old-only classes onto the new primitives, define or delete the 40 undefined ones, wire or delete the 4 orphans, then delete `global.css` so one CSS system remains.
4. **Ship it** ~1h
   `npm run build` + `verify:contract` (72/72), commit, push. Vercel redeploys.
5. **Fix the production detector story**
   Image-only fits Render's RAM where the 1.26 GB audio model does not. Do not let a judge's first click land on an abstention.
6. **Publish one honest number**
   Run `scripts/test_phash.py` and put the recall-versus-transformation table in the README and the deck. One defensible number beats five vague ones.
7. **Quality pass for the 40 percent**
   Reconcile the fusion weights, refresh the README counts, build the Docker image once on a host with a daemon, and label one screen recording as the official demo backup.

ONE OPEN QUESTION, AND IT DECIDES STEPS 5 TO 7
Which event is PRAMAAN actually entered in? The only criteria document we hold is MOSIP Decode 2026 — but its four published problem statements (conformance testing, verifiable-credential interoperability, mobile SSO, face liveness) do not include media forensics, and its timeline (registration closes 14 Sep, judged in November) does not match our 24 Aug Round-1 submission. **Steps 1 to 4 are correct regardless. Steps 5 to 7 need the real rubric and deadline.**